

Hypersolvers: Toward Fast Continuous-Depth Models

Michael Poli^{*,1,3} and Stefano Massaroli^{*,2,3}, Atsushi Yamashita², Hajime Asama², Jinkyoo Park¹

^{*}Equal contribution, ¹KAIST, ²The University of Tokyo, ³DiffEqML



Continuous-Depth Models

We consider the following general Neural ODE formulation:

$$\begin{cases} \dot{\mathbf{z}}(s) = f_{\theta(s)}(s, \mathbf{x}, \mathbf{z}(s)) \\ \mathbf{z}(0) = h_x(\mathbf{x}) \\ \hat{\mathbf{y}}(s) = h_y(\mathbf{z}(s)) \end{cases} \quad s \in \mathcal{S}$$

Neural ODEs are Slow

- **Known:** continuous-depth models are too slow for meaningful **large-scale** or **embedded** applications
- **Existing methods:** regularization terms or augmentation can be used to reduce stiffness and improve inference speed. However, in many settings the vector field $f_{\theta(s)}$ **cannot be modified** e.g control applications.
- We propose a new framework to improve Neural ODE speed in both **inference** and **training**.

A New Paradigm for Neural ODE: Hypersolvers

- **Novel paradigm** to analyze **model-solver interplay**.
- Orthogonal to existing regularization approaches.
- **Idea:** improve pareto efficiency by approximating local residuals with neural networks

General Hypersolver formulation:

$$\mathbf{z}_{k+1} = \mathbf{z}_k + \overbrace{\epsilon \psi(s_k, \mathbf{x}, \mathbf{z}_k)}^{\text{solver step}} + \overbrace{\epsilon^{p+1} g_{\omega}(\epsilon, s_k, \mathbf{x}, \mathbf{z}_k)}^{\text{hypersolver net}} \quad k \in \mathcal{K}$$

$$\mathbf{z}_0 = h_x(\mathbf{x})$$

$$\hat{\mathbf{y}}_k = h_y(\mathbf{z}_k)$$

where ψ is the update step of a p -th order solver

Hypersolver Training Strategies

Assume to have the *exact* solution of the Neural ODE (practically: use adaptive-step solvers with low tolerances). Train g_{ω} with **local truncation residuals** as supervised labels:

$$\mathcal{R}(s_k, \mathbf{z}(s_k), \mathbf{z}(s_{k+1})) = \frac{1}{\epsilon^{p+1}} [\mathbf{z}(s_{k+1}) - \mathbf{z}(s_k) - \epsilon \psi(s_k, \mathbf{x}, \mathbf{z}(s_k))]$$

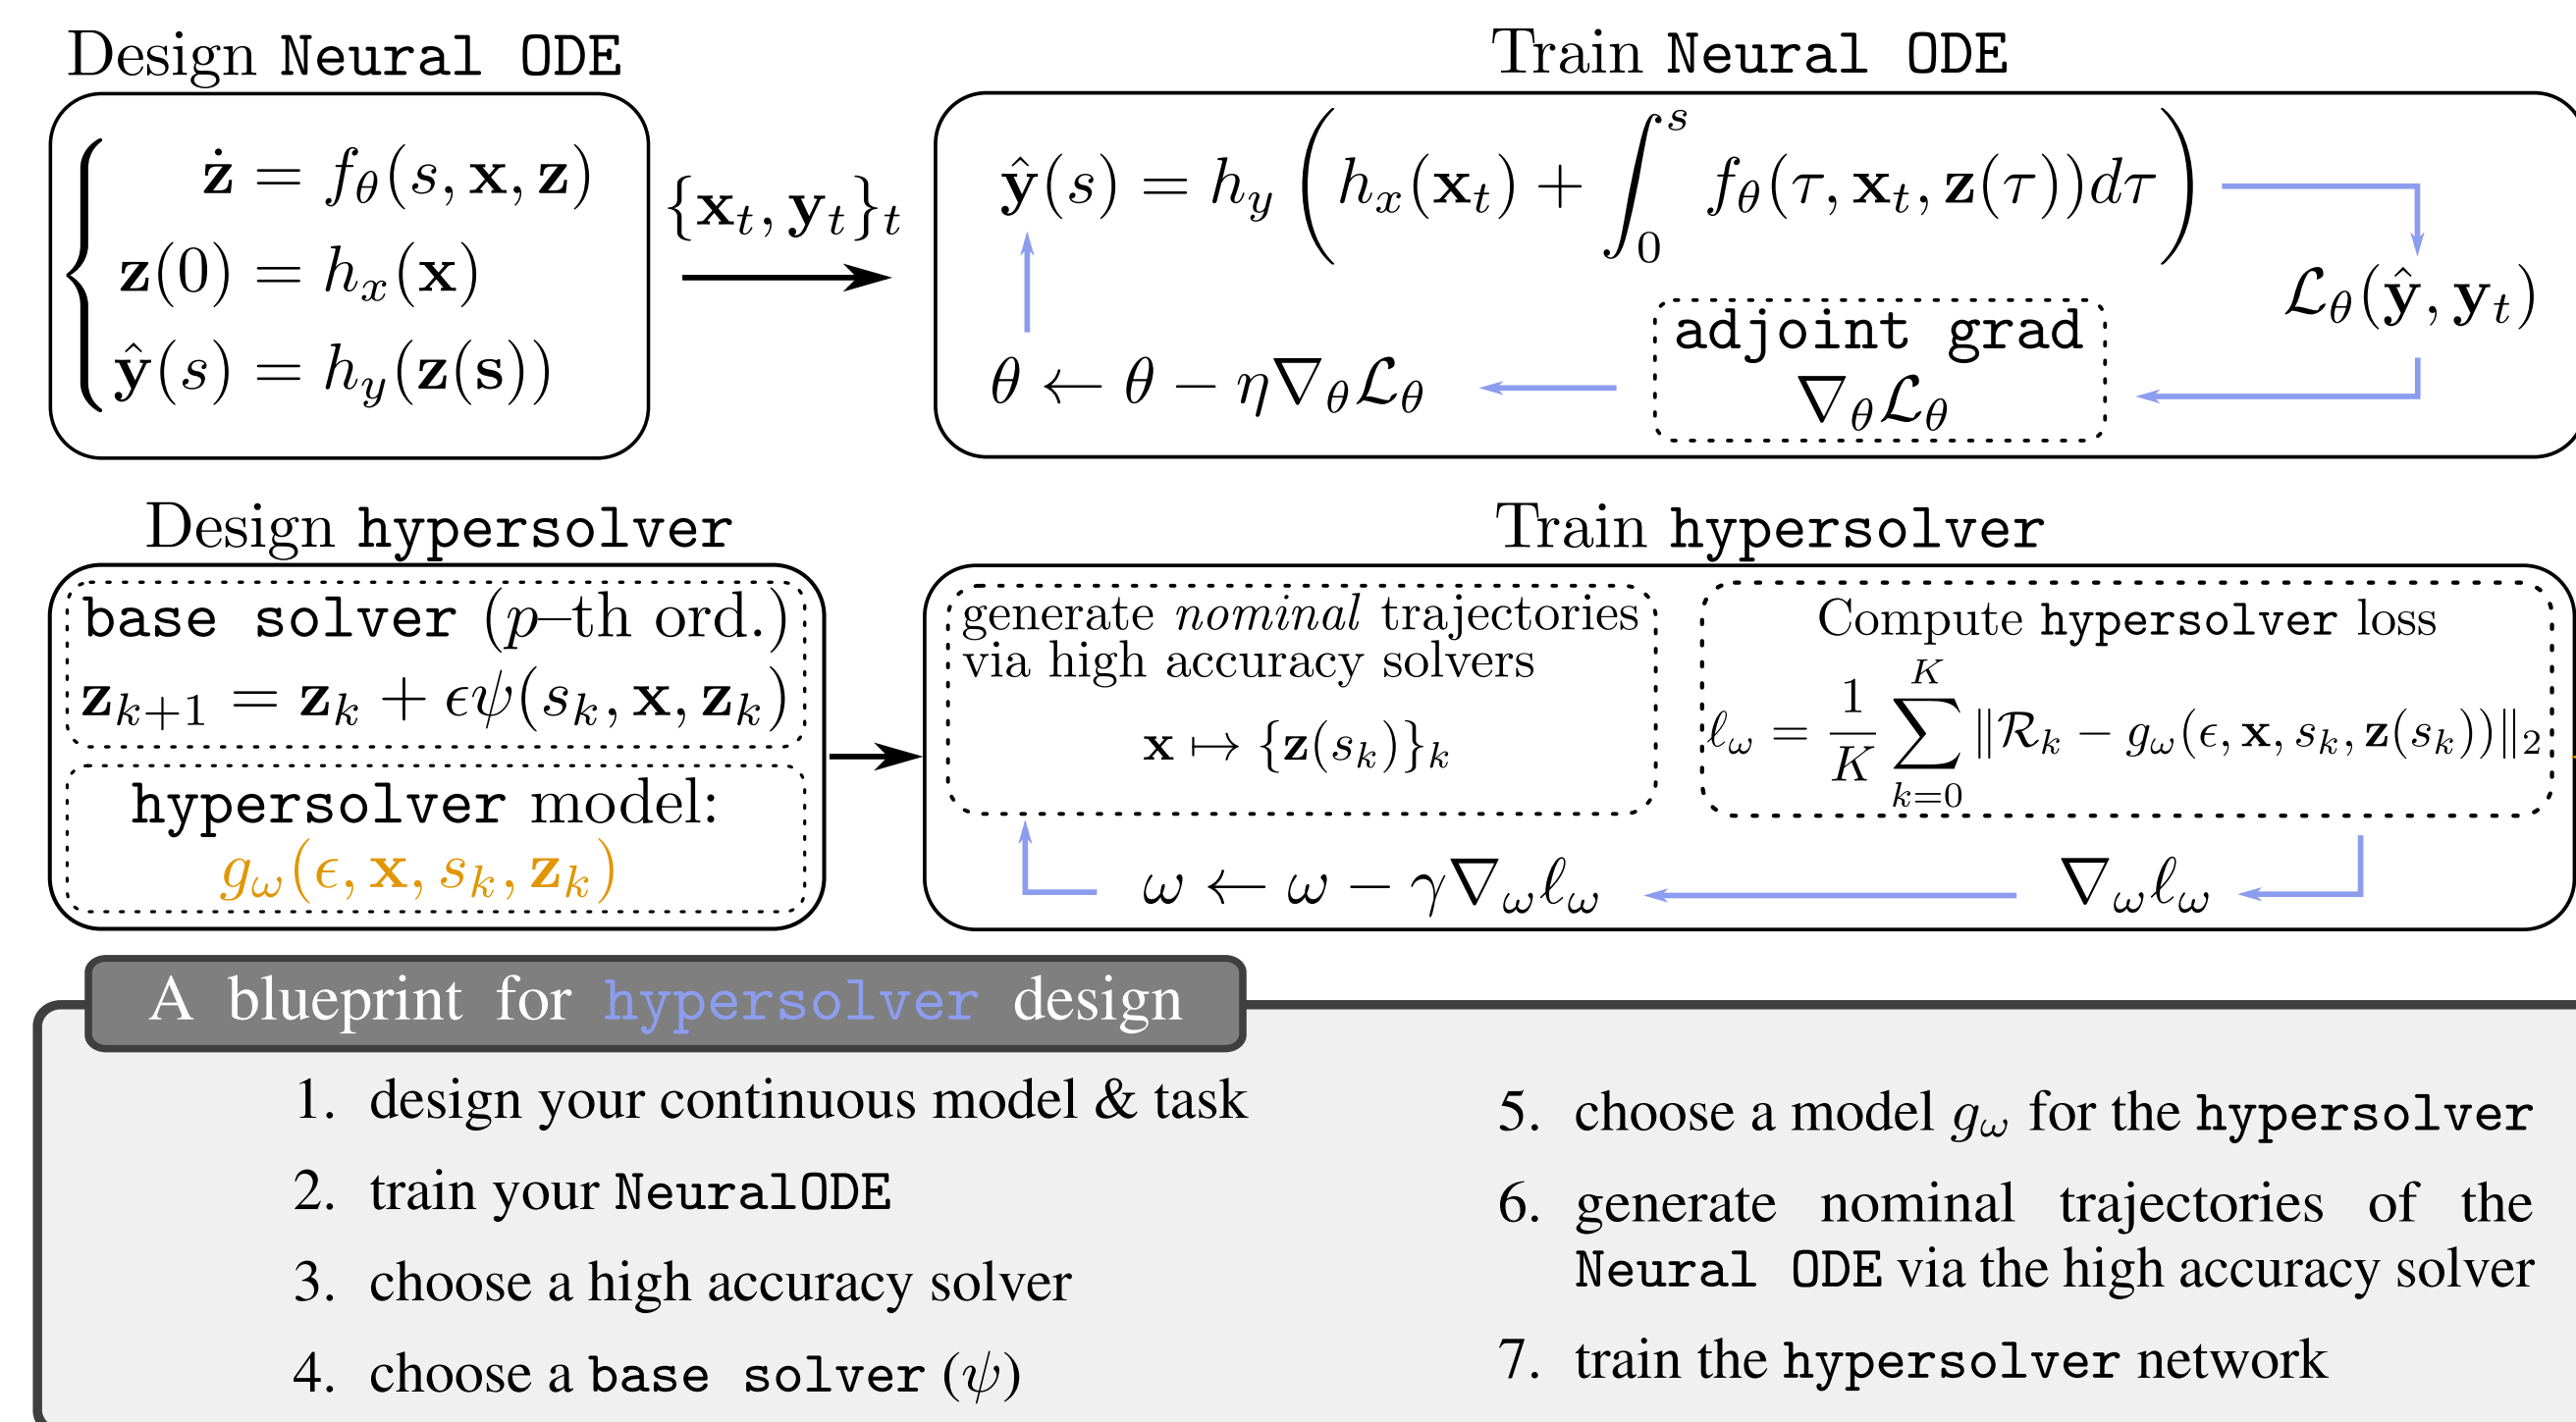
If g_{ω} is a $\mathcal{O}(\delta)$ **approximator** of \mathcal{R} , i.e.

$$\forall k \in \mathbb{N}_{\leq K} \quad \|\mathcal{R}(s_k, \mathbf{z}(s_k), \mathbf{z}(s_{k+1})) - g_{\omega}(\epsilon, s_k, \mathbf{x}, \mathbf{z}(s_k))\|_2 \leq \mathcal{O}(\delta),$$

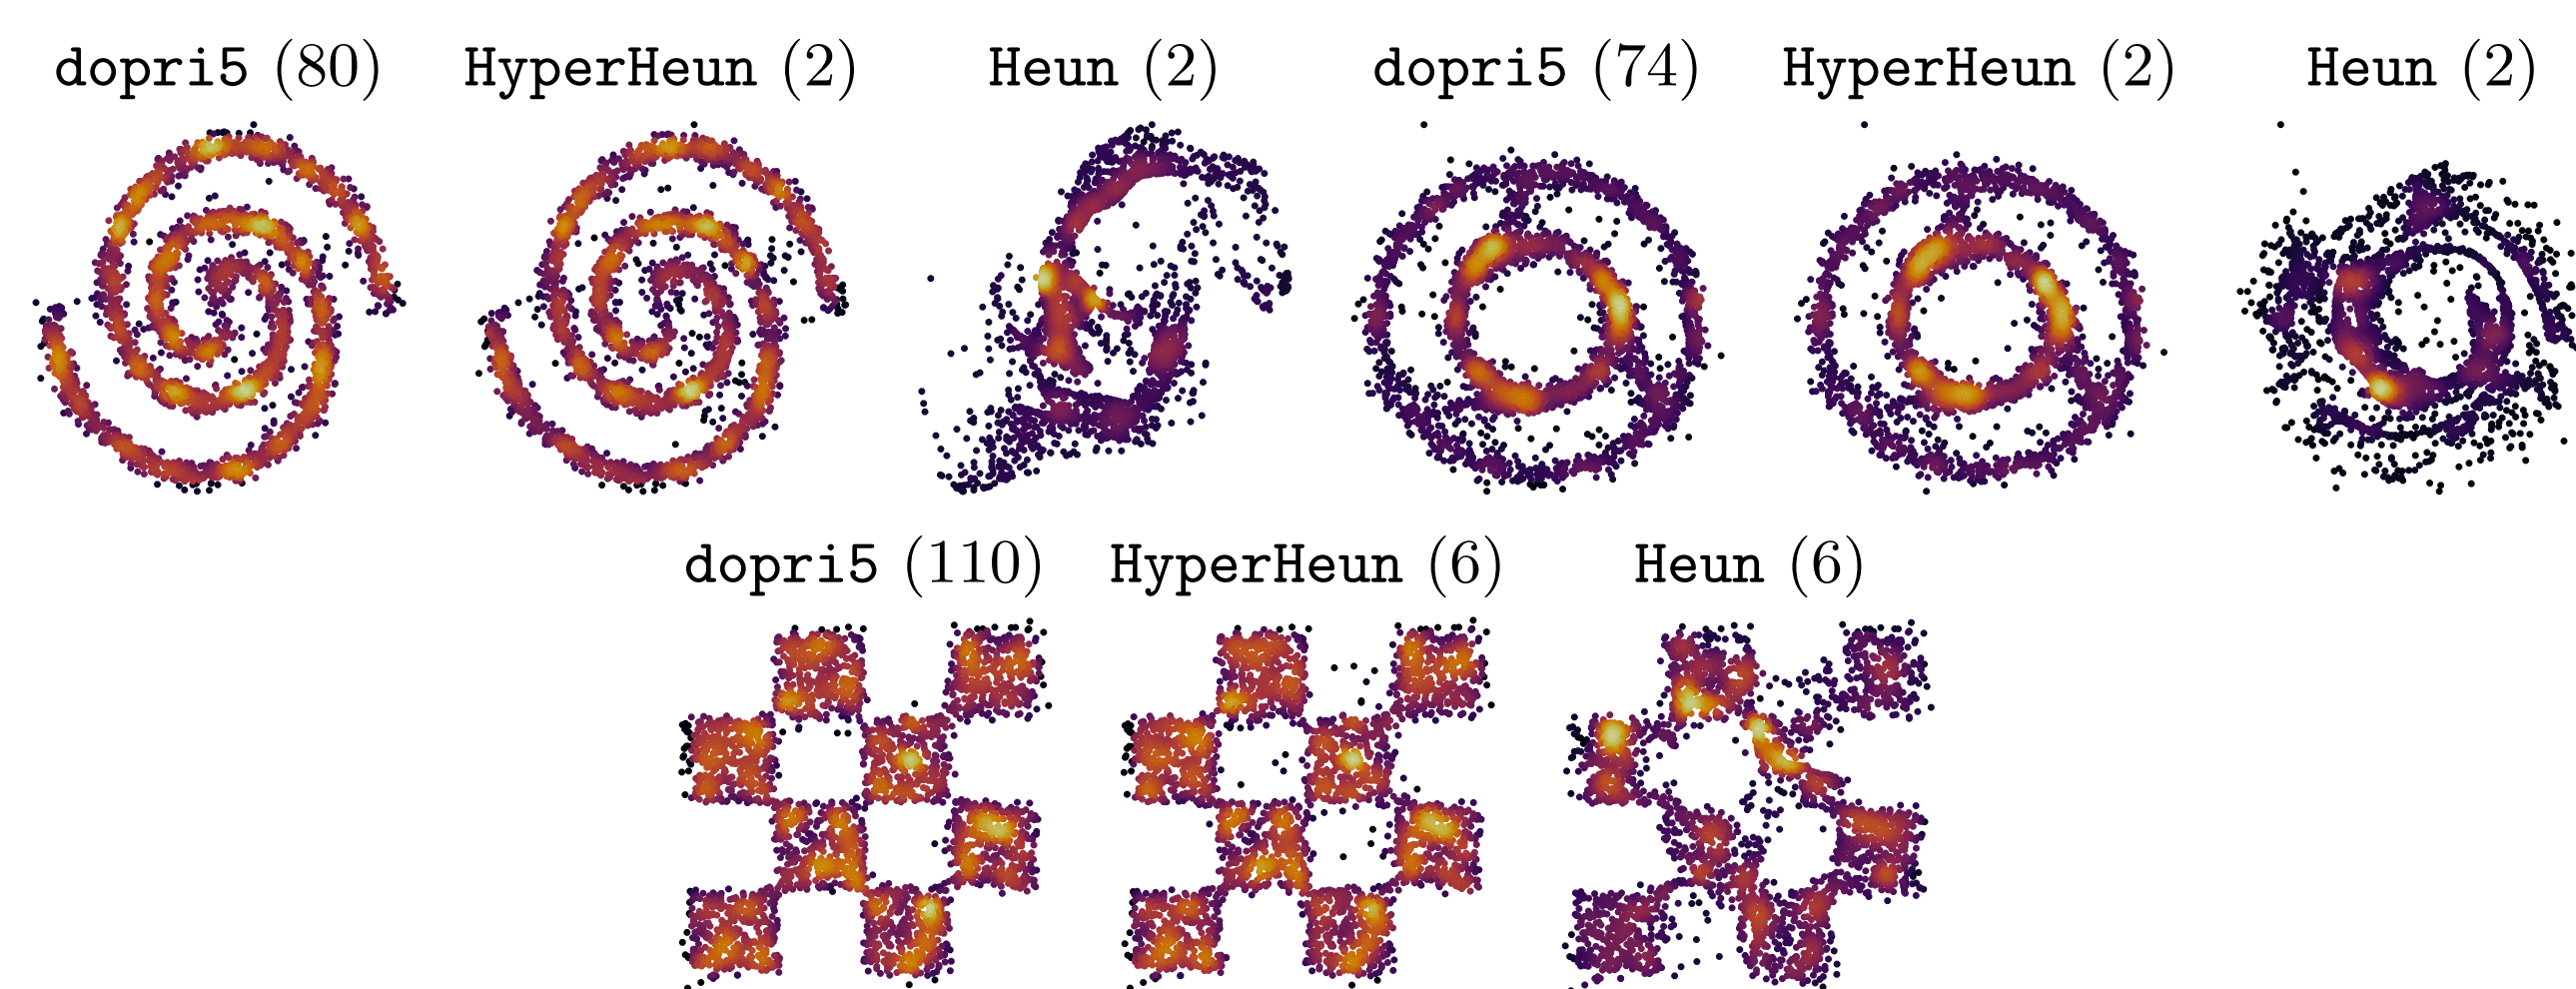
the local truncation error e_k of the hypersolver is $\mathcal{O}(\delta \epsilon^{p+1})$.

Alternative: **trajectory fitting** and **adversarial training**

Blueprint for Hypersolver Design

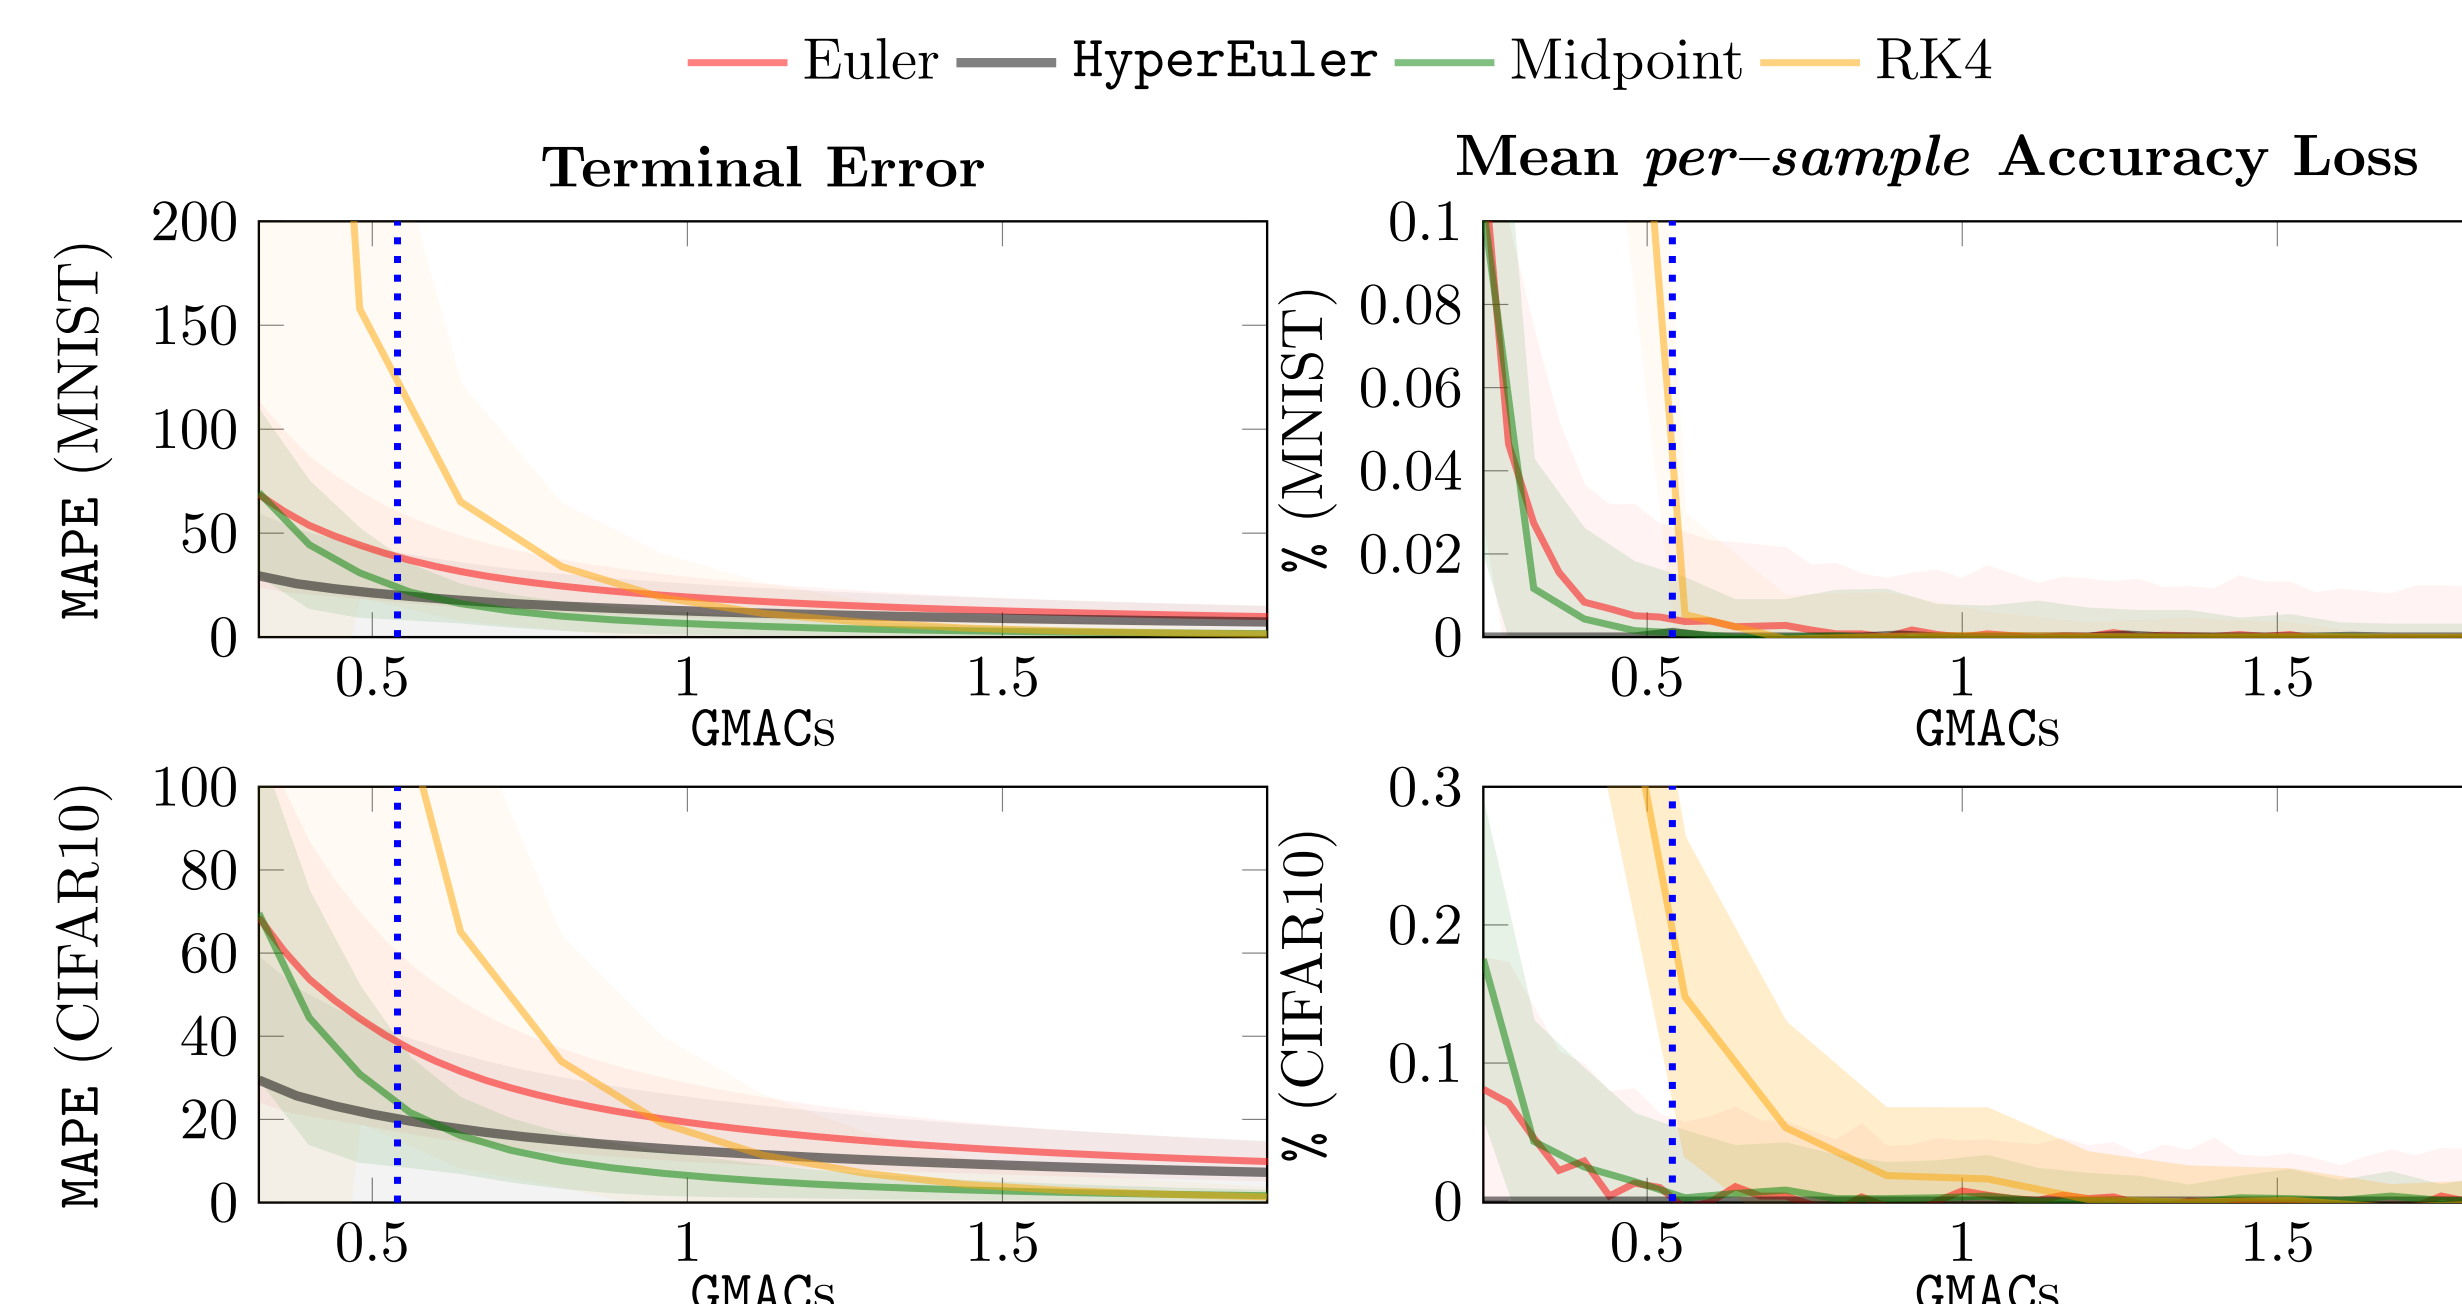


Fast Continuous Normalizing Flow Sampling



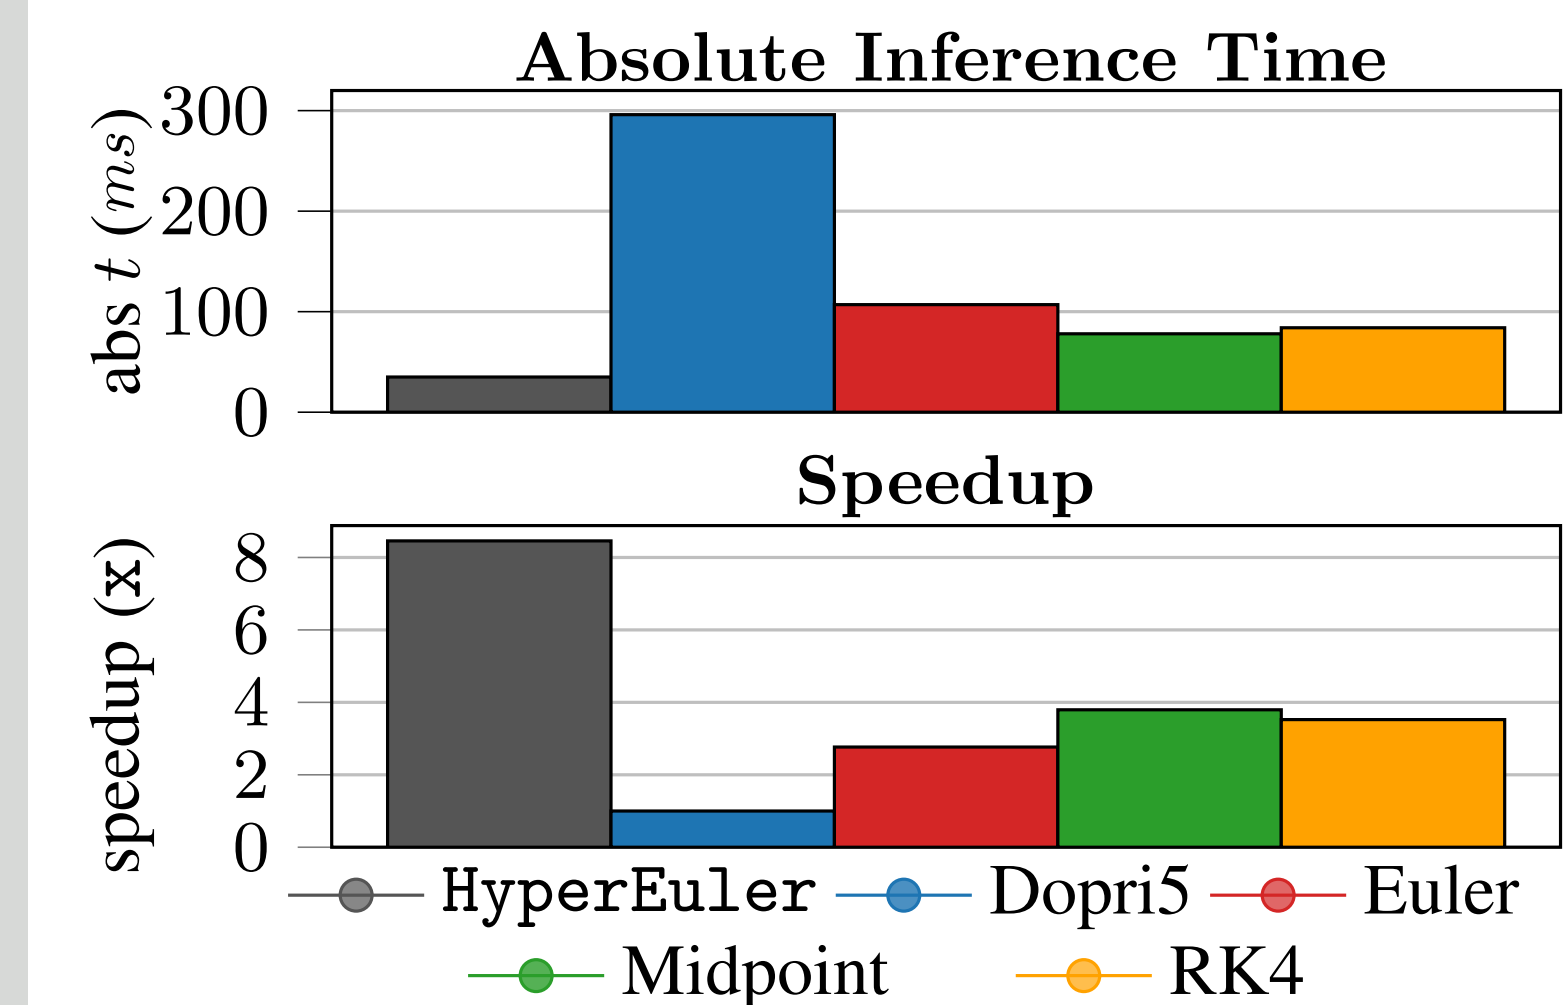
Reconstructed densities (CNFs) solved with various methods. In parenthesis, number of function evaluations. HyperHeun preserves sampled quality with low NFEs.

Pareto Efficiency




Test accuracy loss %-NFE and MAPE-GMAC Pareto fronts of different ODE solvers on MNIST and CIFAR10 test sets. HyperEuler shows higher Pareto efficiency for low function evaluations (NFEs) even over higher-order methods.

Wall-Clock Inference Speedups

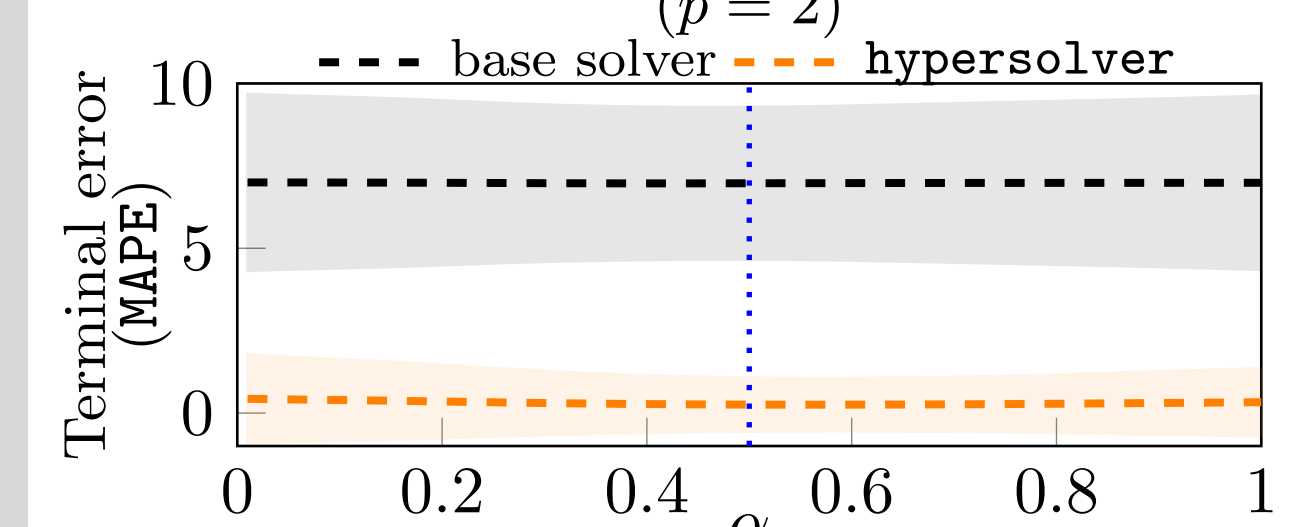


Absolute time (ms) speedup: fixed-step methods over dopri5, MNIST test set. HyperEuler solves the Neural ODE **8x faster** than dopri5 with the **same accuracy**.

Generalization Across Base Solvers



Generalization across base solvers ($p=2$)



Hypersolvers **generalize** across **different base solvers of the same order** (second order explicit solver family parametrized by α).

HyperMidpoint ($\alpha = 0.5$) generalized without finetuning to other base solvers, preserving its Pareto efficiency over the entire α -family.

Hypersolver Overhead

- **Overhead of the method:** evaluation of g_{ω} at each solver step.
- Hypersolver computational overheads decreases as the base solver order p increases. Even in the worst-case scenario (low order), hypersolvers remain Pareto efficient over traditional methods

Accelerating Neural ODE Training

- **Challenge:** ensuring that the hypersolver network remains a $\mathcal{O}(\delta)$ approximator of residuals \mathcal{R} across training iterations.
- Maximizing hypersolver reuse across training iterations represents a path toward faster Neural ODEs (even training).

Beyond Fixed-Step Hypersolvers

- Hypersolver are **not limited** to **fixed-step explicit** solvers.
- Adaptive-step, *predictor-corrector* schemes, different classes of differential equations are also compatible.
- Adversarial training may be used to enhance hypersolver resilience to challenging dynamics.

